

# Controlled Contention Scheduling in Linux Operating System

Presented by  
Phanikar Subodh Krishna Chereddi

# Overview

- Introduction
- CFS scheduler
- Cgroups
- Managing Cgroups
- Performance Measurement Tools
- Results
- Conclusion

# Introduction

- Is scheduling in multicore processor chip a solved problem?
- Modern operating systems like Linux allow both contention and reservation policies.
  - Contention is not good for applications sensitive to interference.
  - Reservation is not good if applications don't use full time interval.
- This is an attempt to implement policies that tradeoff contention scheduling and gang scheduling in Linux operating system using containers.

# Motivation

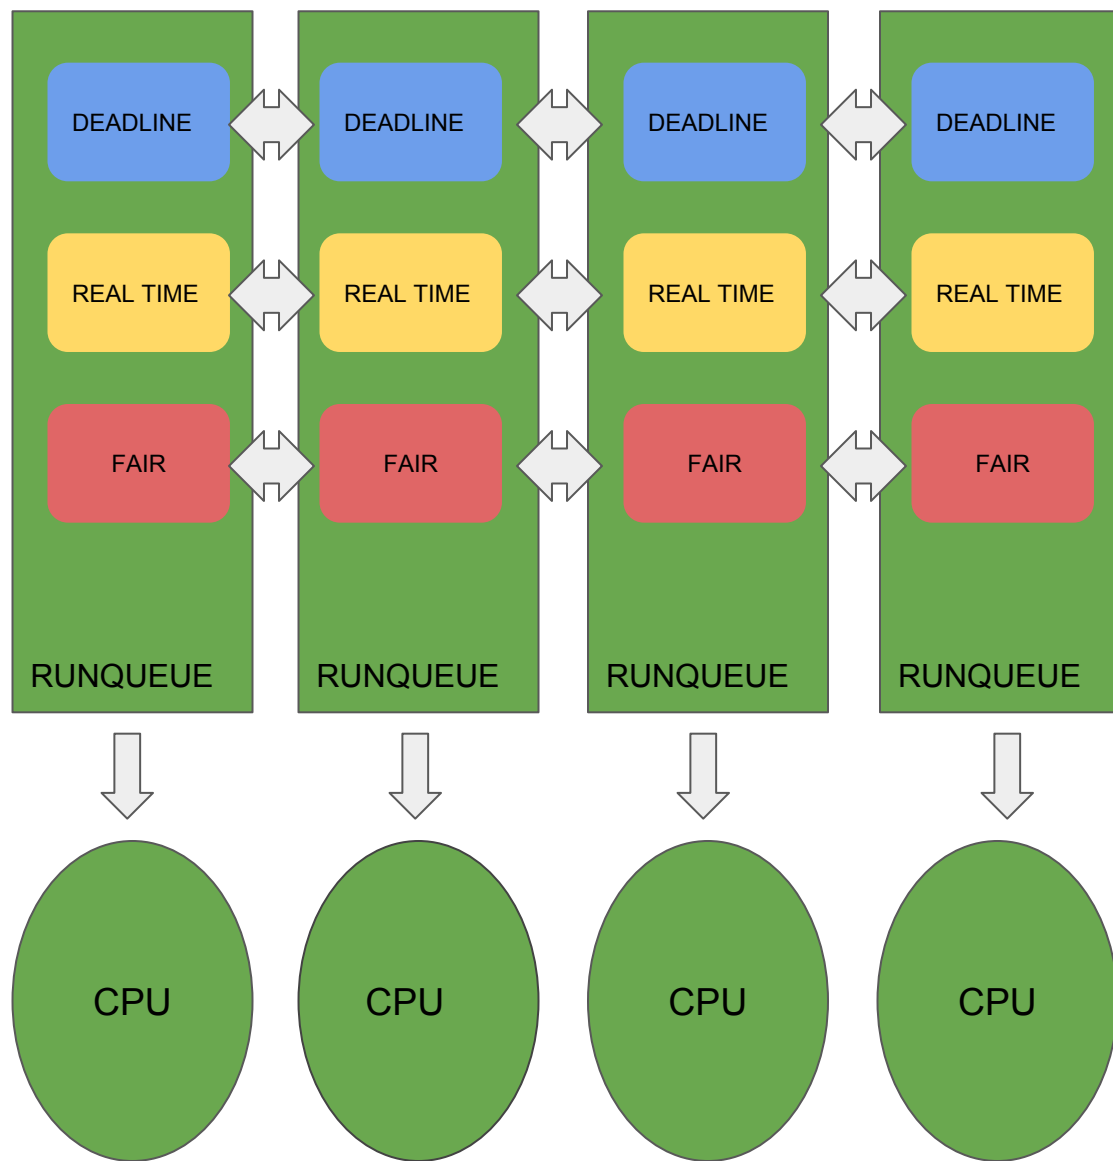
“And you have to realize that there are not very many things that have aged as well as the scheduler. Which is just another proof that scheduling is easy.”

- Linus Torvalds

But still :

- Scheduler is undergoing changes in every new release.
- There is no generic scheduler.
- Especially for desktops booting on Linux, measuring the scheduler performance is difficult.

## Per CPU Runqueues



## Linux Scheduler Architecture

There are 3 scheduling classes in Linux scheduler, their order of priority is:

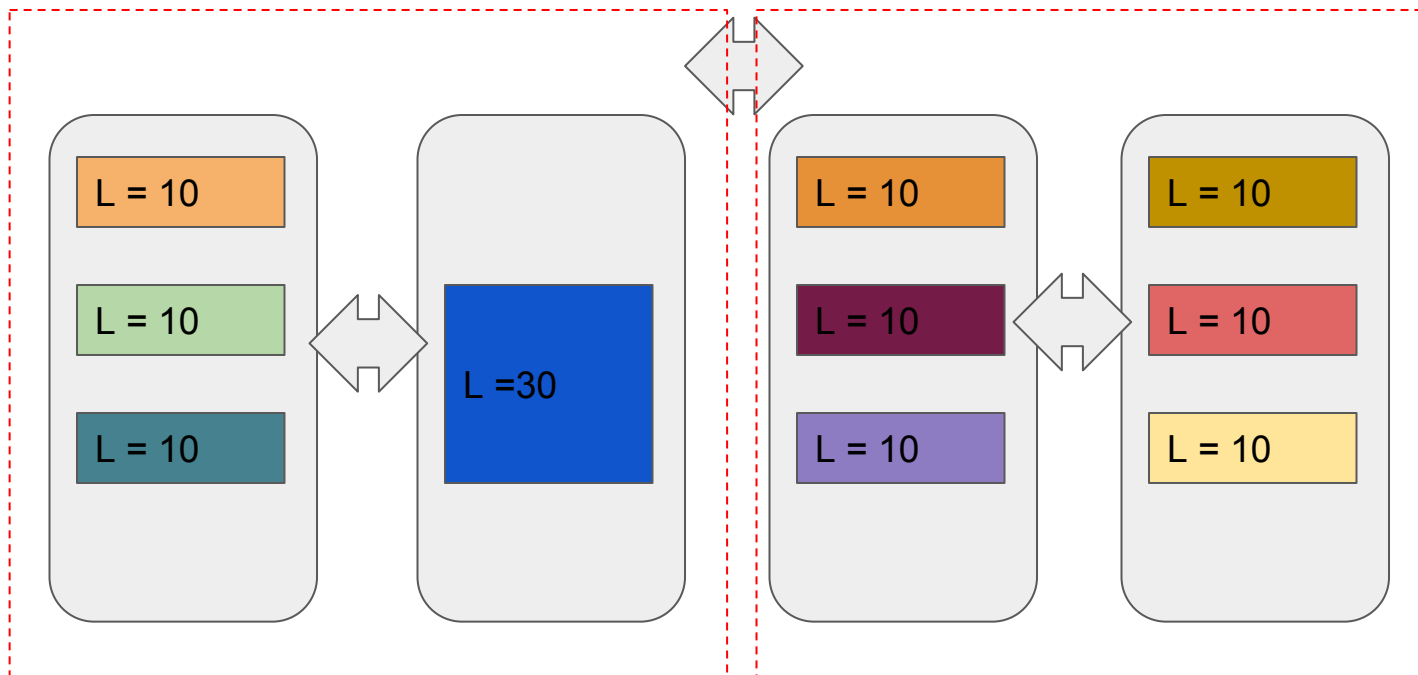
1. Deadline
2. Real time
3. Fair

# Fair Scheduling in Linux kernel

- The virtual runtime of a task is its actual runtime normalized to the total number of running tasks.
- The process in the run queue with shortest virtual runtime has highest priority to be scheduled.
- CFS runqueue is RB-Tree implementation of a priority queue.
- Each node in RB-Tree is a process and they are self-balanced by virtual runtime.
- The leftmost node is one with least virtual runtime so, It has highest priority.


# Load Balancing Among Runqueues

- $\text{Load} = (\text{Weight of process}) * (\% \text{ Utilization of CPU})$
- Load average is metric used to balance the runnable processes among CPUs.



# Gang Scheduling


- Strict gang scheduling does not allow threads of other application to be scheduled, may lead to wasted computational power.






# Gang Scheduling

- Allowing other application threads to run when there is an ideal CPU may cause interference.



# Pinning Application

- Pining guarantees that application will run only in that CPU.
- It does not mean that only that application will be running exclusively.
- The kernel threads will also be scheduled on all the CPUs.



# Containers in Linux Kernel

- Container types:
  - Cgroups
  - Namespaces
- Containers are lightweight virtual machines:
  - Own process space, own network interface, can run stuff as root, can install packages, can run services.
- But a little different from virtual machines:
  - Uses the host kernel, can't boot a different OS, can't have its own modules, doesn't need init as PID 1.

# Cgroups

- **blkio** — this subsystem sets limits on input/output access to and from block devices such as physical drives (disk, solid state, or USB).
- **cpu** — this subsystem uses the scheduler to provide cgroup tasks access to the CPU.
- **cpuacct** — this subsystem generates automatic reports on CPU resources used by tasks in a cgroup.
- **cpuset** — this subsystem assigns individual CPUs (on a multicore system) and memory nodes to tasks in a cgroup.
- **devices** — this subsystem allows or denies access to devices by tasks in a cgroup.

# Cgroups

- **freezer** — this subsystem suspends or resumes tasks in a cgroup.
- **memory** — this subsystem sets limits on memory use by tasks in a cgroup and generates automatic reports on memory resources used by those tasks.
- **net\_cls** — this subsystem tags network packets with a class identifier (classid) that allows the Linux traffic controller (tc) to identify packets originating from a particular cgroup task.
- **net\_prio** — this subsystem provides a way to dynamically set the priority of network traffic per network interface.
- **ns** — the *namespace* subsystem.

# Cgroups

```
phanikar@phanikar-XPS-8900:~$ ls -l /sys/fs/cgroup/
total 0
dr-xr-xr-x  2 root root  0 Dec  6 10:22 blkio
drwxr-xr-x  2 root root 60 Dec  6 10:22 cgmanager
lrwxrwxrwx  1 root root 11 Dec  6 10:22 cpu -> cpu,cpuacct
lrwxrwxrwx  1 root root 11 Dec  6 10:22 cpuacct -> cpu,cpuacct
dr-xr-xr-x  2 root root  0 Dec  6 10:22 cpu,cpuacct
dr-xr-xr-x 10 root root  0 Dec  6 10:22 cpuset
dr-xr-xr-x  5 root root  0 Dec  6 10:22 devices
dr-xr-xr-x 10 root root  0 Dec  6 10:22 freezer
dr-xr-xr-x  2 root root  0 Dec  6 10:22 hugetlb
dr-xr-xr-x  2 root root  0 Dec  6 10:22 memory
lrwxrwxrwx  1 root root 16 Dec  6 10:22 net_cls -> net_cls,net_prio
dr-xr-xr-x  2 root root  0 Dec  6 10:22 net_cls,net_prio
lrwxrwxrwx  1 root root 16 Dec  6 10:22 net_prio -> net_cls,net_prio
dr-xr-xr-x  2 root root  0 Dec  6 10:22 perf_event
dr-xr-xr-x  5 root root  0 Dec  6 10:22 pids
dr-xr-xr-x  5 root root  0 Dec  6 10:22 systemd
```

# Freezer Cgroup

- Allows to **freeze / thaw** a group of processes or threads.
- Similar in functionality to mass **SIGSTOP/SIGCONT!!!**
- Cannot be detected by the processes (unlike SIGSTOP/SIGCONT).
- Doesn't impede ptrace/debugging.
- Specific use cases:
  - Cluster batch scheduling
  - Process migration

# Freezer Cgroup

- **freezer.state:** Read-write.  
Shows state as THAWED / FROZEN / FREEZING.
- **freezer.self\_freezing:** Read only.  
Shows the self-state. 0 if the self-state is “THAWED”; otherwise, 1. This value is 1 iff the last write to freezer.state was “FROZEN”.
- **freezer.parent\_freezing:** Read only.  
Shows the parent-state. 0 if none of the cgroup's ancestors is frozen; otherwise, 1.
- **NOTE:** The root cgroup is non-freezable and the above interface files don't exist.



# Cpuset Cgroup

- Pin groups to specific CPU(s)
- Reserve CPUs for specific apps
- Avoid processes bouncing between CPUs
- Also relevant for NUMA systems
- Provides extra dials and knobs
  - Per zone memory pressure, process migration costs...

# Cpuset Cgroup

- **cpuset.cpus**  
List of the physical numbers of the CPUs on which processes in that cpuset are allowed to execute.
- **cpuset.cpu\_exclusive** Flag (0 or 1)  
If set (1), the cpuset has exclusive use of its CPUs (no sibling or cousin cpuset may overlap CPUs).
- **cpuset.sched\_load\_balance** Flag (0 or 1)  
If set (1, the default) the kernel will automatically load balance processes in that cpuset.
- **cpuset.sched\_relax\_domain\_level** - For NUMA machine
  - Integer, between -1 and a small positive value.
  - The *sched\_relax\_domain\_level* controls the width of the range of CPUs over which the kernel scheduler performs immediate rebalancing.

# Cpu,cpuacct Cgroup

- Keeps track of user/system CPU time
- Keeps track of usage per CPU
- Allows to set weights
  - `cpu.shares`
- How regularly a cgroup's access to CPU resources should be reallocated?
  - `cpu.cfs_period_us`
- Specifies the total amount of time in which all tasks in a cgroup can run during one period
  - `cpu.cfs_quota_us`

# Managing Cgroups

- Cgroup pseudo filesystem is mounted in **/sys/fs/cgroup/**
- Commands:
  - **cgcreate** -t uid:gid -a uid:gid -g subsystems:path
  - **cgdelete** subsystems:path
  - **cgset** -r parameter=value path\_to\_cgroup
  - **cgexec** -g subsystems:path\_to\_cgroup command arguments
  - **lscgroup**

# Managing Cgroups

- Manually pushing a PID to a Cgroup
  - `echo pid > /sys/fs/cgroup/{subsystem}/{path-to-hierarchy}/tasks`
- Writing parameter to configuration files
  - `echo THAWED > /sys/fs/cgroup/freezer/app1/freezer.state`

# Scheduler Analysis / Job Management tools

- Htop
  - Colorful and Interactive
  - Can view resource utilization per process

```
phanikar@phanikar-XPS-8900: ~/Desktop/parsec-2.1/pkgs/apps/blackscholes/inst/amd64-linux.gcc.pre/bin 231x33

1 [|||||] 93.4% 5 [|||||] 100.0%
2 [|||||] 93.4% 6 [|||||] 100.0%
3 [|||||] 77.2% 7 [|||||] 100.0%
4 [|||||] 97.4% 8 [|||||] 100.0%
Mem[|||||] 5.14G/15.6G Tasks: 136, 327 thr: 19 running
Swp[|||||] 2.39M/15.9G Load average: 2.77 0.81 0.54
Uptime: 1 day, 01:25:41

PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
20917 phanikar 20 0 3315M 1756M 146M S 31.8 11.0 32:09.49 /usr/lib/firefox/firefox
1088 root 20 0 371M 100M 77712 S 0.0 0.6 22:40.05 /usr/lib/xorg/Xorg -core :0 -seat seat0 -auth /var/run/lightdm/root/:0 -nolisten tcp vt7 -novtswitch
3298 phanikar 20 0 1549M 239M 65016 S 0.0 1.5 8:12.77 compiz
3175 phanikar 20 0 33836 4680 3172 R 0.7 0.0 5:38.54 htop
10769 phanikar 20 0 770M 61292 33616 S 0.0 0.4 1:52.03 /usr/bin/python /usr/bin/terminator
20957 phanikar 20 0 3315M 1756M 146M S 0.0 11.0 0:39.71 /usr/lib/firefox/firefox
20944 phanikar 20 0 3315M 1756M 146M R 0.0 11.0 0:18.82 /usr/lib/firefox/firefox
20966 phanikar 20 0 3315M 1756M 146M S 0.0 11.0 0:14.14 /usr/lib/firefox/firefox
3111 phanikar 20 0 471M 35744 23084 S 0.0 0.2 0:17.34 /usr/lib/ibus/ibus-ui-gtk3
3093 phanikar 20 0 426M 20672 5412 S 0.0 0.1 0:31.82 /usr/bin/ibus-daemon --daemonize --xim
3252 phanikar 20 0 512M 29596 20996 S 0.0 0.2 0:10.30 /usr/lib/x86_64-linux-gnu/bamf/bamfdaemon
3026 phanikar 20 0 44048 4456 2816 S 0.0 0.0 0:07.41 dbus-daemon --fork --session --address=unix:abstract=/tmp/dbus-RZHMdS5EbQ
3415 phanikar 20 0 1549M 239M 65016 S 0.0 1.5 0:04.31 compiz
10773 phanikar 20 0 770M 61292 33616 S 0.0 0.4 0:01.50 /usr/bin/python /usr/bin/terminator
3467 phanikar 20 0 984M 33236 26432 S 0.0 0.2 0:30.00 nm-applet
899 messagebu 20 0 44376 5148 3364 S 0.0 0.0 0:22.27 /usr/bin/dbus-daemon --system --address=systemd: --nofork --nopidfile --systemd-activation
951 root 20 0 590M 16980 13396 S 0.0 0.1 0:29.89 /usr/sbin/NetworkManager --no-daemon
20933 phanikar 20 0 3315M 1756M 146M S 0.0 11.0 0:11.56 /usr/lib/firefox/firefox
20982 phanikar 20 0 3315M 1756M 146M S 0.0 11.0 0:03.94 /usr/lib/firefox/firefox
1060 root 20 0 590M 16980 13396 S 0.0 0.1 0:10.60 /usr/sbin/NetworkManager --no-daemon
3303 phanikar 20 0 635M 47836 24072 S 0.0 0.3 0:20.57 /usr/lib/x86_64-linux-gnu/unity/unity-panel-service
3086 phanikar 20 0 39864 304 0 S 0.0 0.0 0:02.43 upstart-dbus-bridge --daemon --system --user --bus-name system

F1 Help F2 Setup F3 Search F4 Filter F5 Tree F6 Sort By F7 Nice F8 Nicer F9 Kill F10 Quit
```

# Scheduler Analysis / Job Management tools

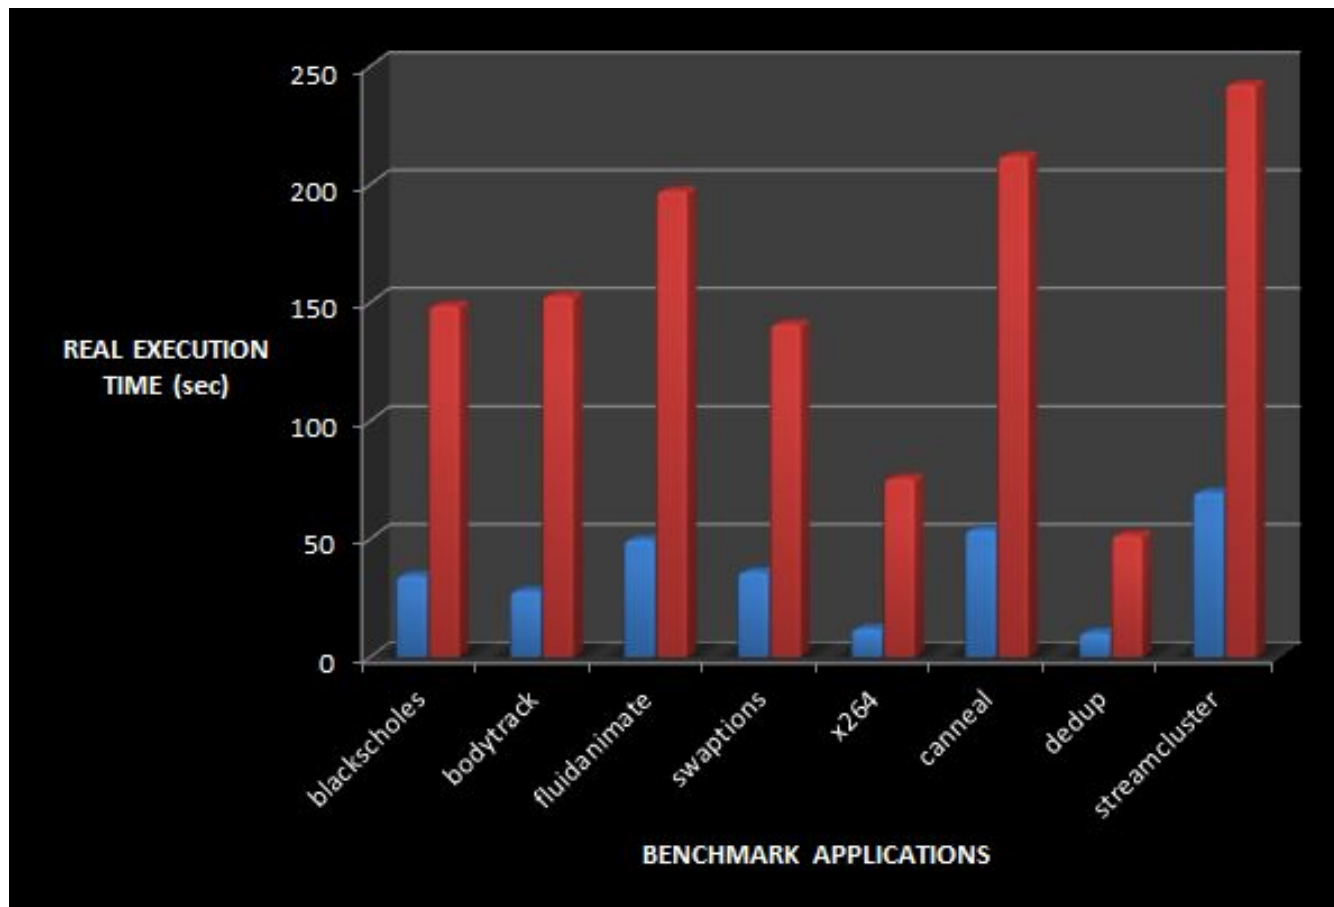
- Berkeley Packet Filter (BPF)
  - Runqlat, Cpudist are part of BPF to measure runqueue latency and cpu distribution
  - Easy to debug, written in Python

```
root@phanikar-XPS-8900:/usr/share/bcc/tools# ./cpudist
Tracing on-CPU time... Hit Ctrl-C to end.
^C
  usecs      : count      distribution
    0 -> 1    : 161      ***
    2 -> 3    : 53       *
    4 -> 7    : 68       *
    8 -> 15   : 66       *
   16 -> 31   : 65       *
   32 -> 63   : 160      ***
   64 -> 127  : 147      ***
  128 -> 255  : 211      ****
  256 -> 511  : 275      *****
  512 -> 1023 : 554      *****
 1024 -> 2047 : 388      *****
 2048 -> 4095 : 1712     *****
 4096 -> 8191 : 183      ****
 8192 -> 16383: 215      *****
16384 -> 32767: 103      **
32768 -> 65535: 83       *
65536 -> 131071: 18
```

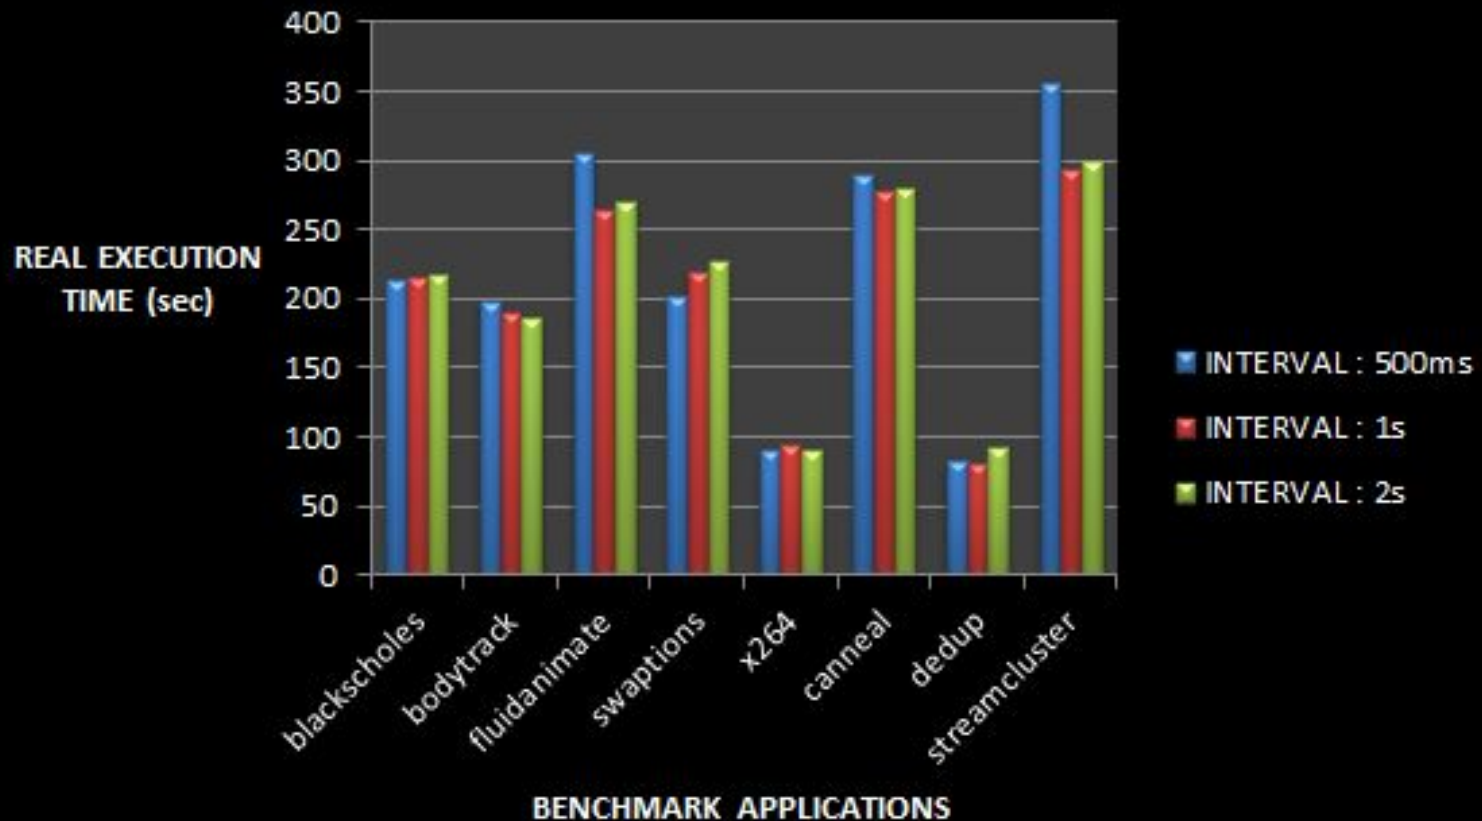
# Experiment Results




# Runtime of Individual Application VS Runtime of Applications Running Parallely Using CFS



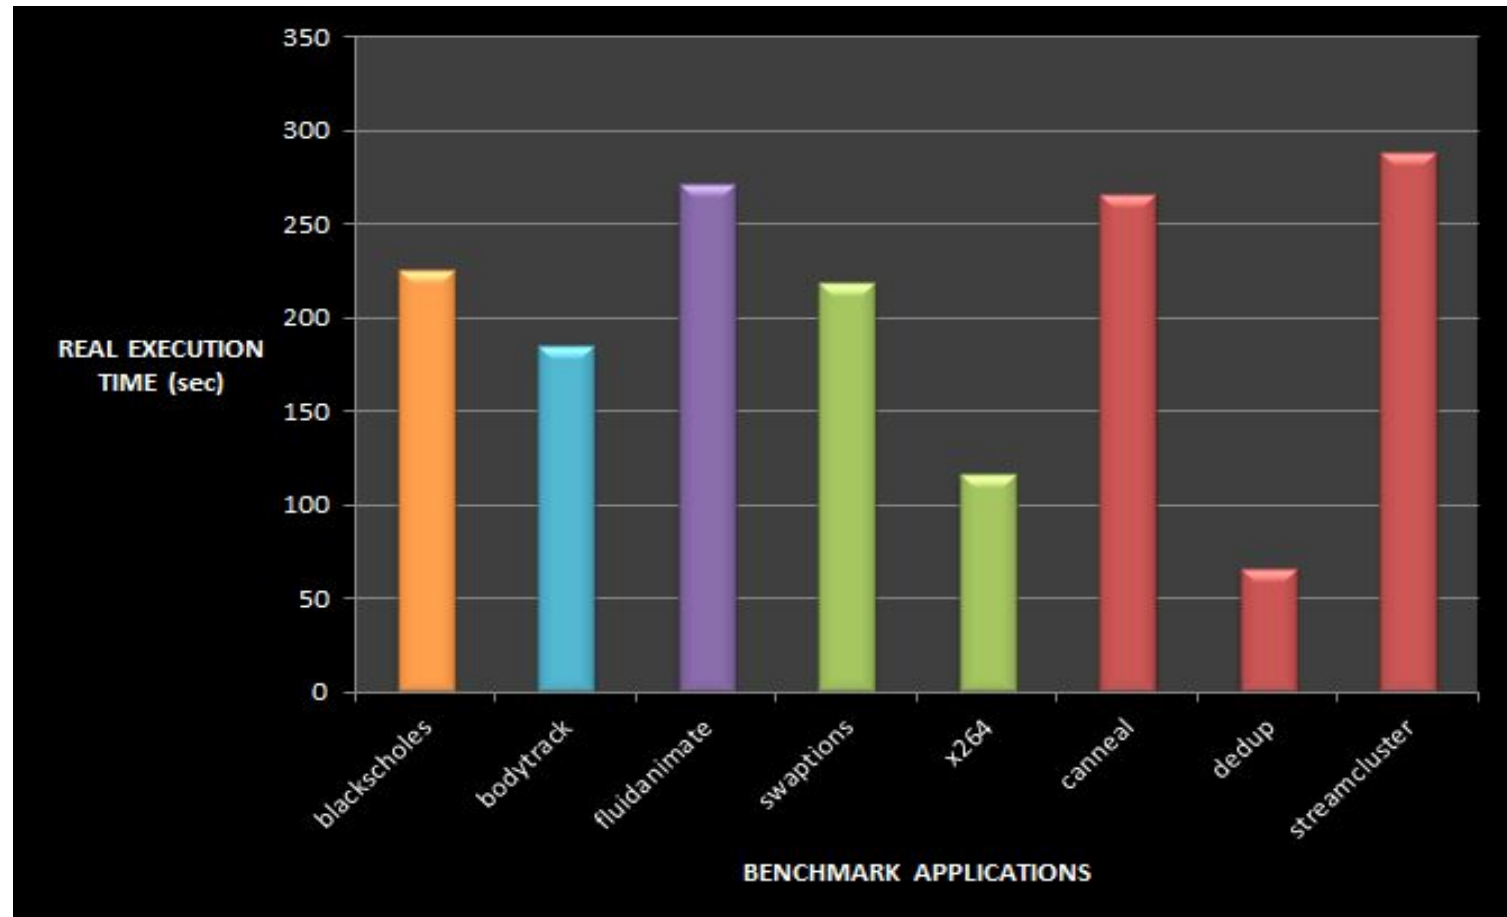
# Runtime of Applications Using Gang Scheduling With Different Switching Intervals




# Runtime of Applications Using Pinning With Different Thread Counts



# Runtime of Applications When Co Scheduled



# Runtime of Applications When Co Scheduled



# Conclusions & Future Work

- Terminology used in the kernel leads to confusion often.
- Available books are based on older versions but, there is minor release almost every month.
- There was unexpected behaviour as policy was implemented in userspace instead of kernel space.
- I wish to implement Contention Controlled scheduling policies in the Linux Scheduler in the future.

# References

1. [Controlled Contention: Balancing Contention and Reservation in Multicore Application Scheduling](#)
2. [The Linux Scheduler: a Decade of Wasted Cores](#)
3. [extended Berkeley Packet Filters](#)
4. [Cgroups](#)

**THANK YOU**



# Running Applications

RUNNING INDIVIDUAL APPLICATION IN SEQUENCE			
application	real time (sec)	virtual time in userspace (sec)	threads
blackscholes	34.345	180.756	8
bodytrack	27.86	178.192	8
fluidanimate	49.598	379.628	8
swaptions	35.933	281.36	8
x264	12.03	87.732	8
canneal	53.803	186.992	8
dedup	10.269	29.724	8
streamcluster	69.606	545.32	8
<b>total</b>	<b>293.444</b>	1869.704	

RUNNING ALL APPLICATIONS IN PARALLEL			
application	real time (sec)	virtual time in userspace (sec)	threads
blackscholes	148.689	188.076	8
bodytrack	152.754	176.084	8
fluidanimate	197.296	381.064	8
swaptions	140.818	288.36	8
x264	75.535	87.824	8
canneal	212.055	195.296	8
dedup	51.577	29.084	8
streamcluster	<b>242.338</b>	498.644	8
<b>total</b>	<b>242.345</b>	1844.432	

RUNNING ALL APPLICATION USING FREEZER			
SWITCHING INTERVAL : 500ms			
application	real time (sec)	virtual time in userspace (sec)	threads
blackscholes	212.874	181.192	8
bodytrack	197.814	179.36	8
fluidanimate	304.145	436.276	8
swaptions	201.32	281.54	8
x264	90.243	88.712	8
canneal	288.907	187.028	8
dedup	82.996	29.512	8
streamcluster	<b>355.782</b>	701.66	8
<b>total</b>	<b>355.782</b>	2085.28	

RUNNING ALL APPLICATION USING FREEZER			
SWITCHING INTERVAL : 1s			
application	real time (sec)	virtual time in userspace (sec)	threads
blackscholes	216.118	180.432	8
bodytrack	190.264	176.98	8
fluidanimate	264.042	379.636	8
swaptions	218.217	281.4	8
x264	93.331	87.676	8
canneal	276.784	186.756	8
dedup	80.047	29.856	8
streamcluster	<b>293.788</b>	545.18	8
<b>total</b>	<b>293.869</b>	1868.016	

RUNNING ALL APPLICATION USING FREEZER			
SWITCHING INTERVAL : 2s			
application	real time (sec)	virtual time in userspace (sec)	threads
blackscholes	216.578	181.38	8
bodytrack	185.661	176.728	8
fluidanimate	269.278	379.888	8
swaptions	226.851	281.832	8
x264	90.56	87.888	8
canneal	280.399	186.928	8
dedup	92.934	29.932	8
streamcluster	<b>298.522</b>	545.348	8
<b>total</b>	<b>299.399</b>	1869.96	

application	real time (sec)	virtual time in userspace (sec)	threads	cpu usage
blackscholes	225.09	188.62	8	83%CPU
bodytrack	185.31	175.74	8	95%CPU
fluidanimate	270.97	391.72	8	144%CPU
swaptions	218.69	290.05	8	132%CPU
x264	115.8	88.3	8	76%CPU
canneal	265.99	218.17	8	82%CPU
dedup	65.69	28.96	8	49%CPU
streamcluster	288.34	564.4	8	196%CPU
<b>total</b>	<b>288.34</b>	1945.96		

application	real time (sec)	virtual time in userspace (sec)	threads	cpu usage
blackscholes	162.56	189.66	8	116%CPU
bodytrack	185.33	176.22	8	95%CPU
fluidanimate	236.25	376.8	8	159%CPU
swaptions	150.92	292.68	8	193%CPU
x264	64.92	86.98	8	134%CPU
canneal	241.95	222.14	8	92%CPU
dedup	101.09	28.88	8	32%CPU
streamcluster	268.8	538.13	8	200%CPU
<b>total</b>	<b>268.8</b>	1911.49		